

# Review Model\_evaluation\_classification

COSC 3337: Data Science I  
Instructor: Nouhad Rizk

## 1 Model Evaluation 2 – Confidence Intervals and Bootstrapping

```
[2]: %matplotlib inline
import numpy as np
import matplotlib.pyplot as plt
```

## 2 Iris Feature by Class Distribution in Random Subsampling

```
[3]: import matplotlib.pyplot as plt
import numpy as np
import pandas as pd
from mlxtend.data import iris_data
from mlxtend.preprocessing import shuffle_arrays_unison

X, y = iris_data()

X, y = shuffle_arrays_unison([X, y], random_seed=123)
X_train, X_test = X[:100], X[100:150]
y_train, y_test = y[:100], y[100:150]
```

```
[4]: np.bincount(y)
```

```
[4]: array([50, 50, 50])
```

```
[5]: np.bincount(y_train)
```

```
[5]: array([38, 28, 34])
```

```
[6]: np.bincount(y_test)
```

```
[6]: array([12, 22, 16])
```

```

[7]: df = pd.DataFrame(X)
      df['class'] = y

      df_train = pd.DataFrame(X_train)
      df_train['class'] = y_train

      df_test = pd.DataFrame(X_test)
      df_test['class'] = y_test

[8]: def stackhist(x, y, **kws):
      grouped = x.groupby(y)
      data = [d for _, d in grouped]
      labels = [l for l, _ in grouped]
      plt.hist(data,
                histtype="barstacked",
                label=labels,
                alpha=0.8,
                density=True,
                bins=np.arange(4.0, 8.1, 0.25))
      plt.ylim([0, 0.8])
      plt.xlim([4, 8])
      plt.xlabel('Sepal Width [cm]')
      plt.ylabel('Frequency')


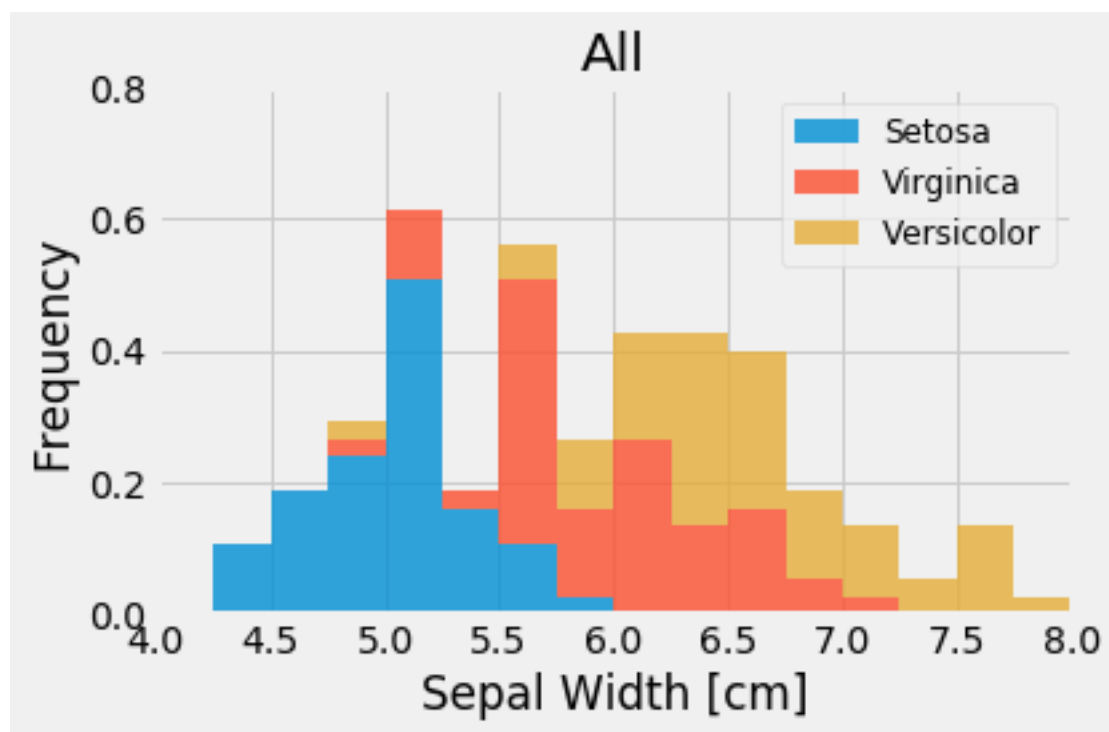
      with plt.style.context('fivethirtyeight'):

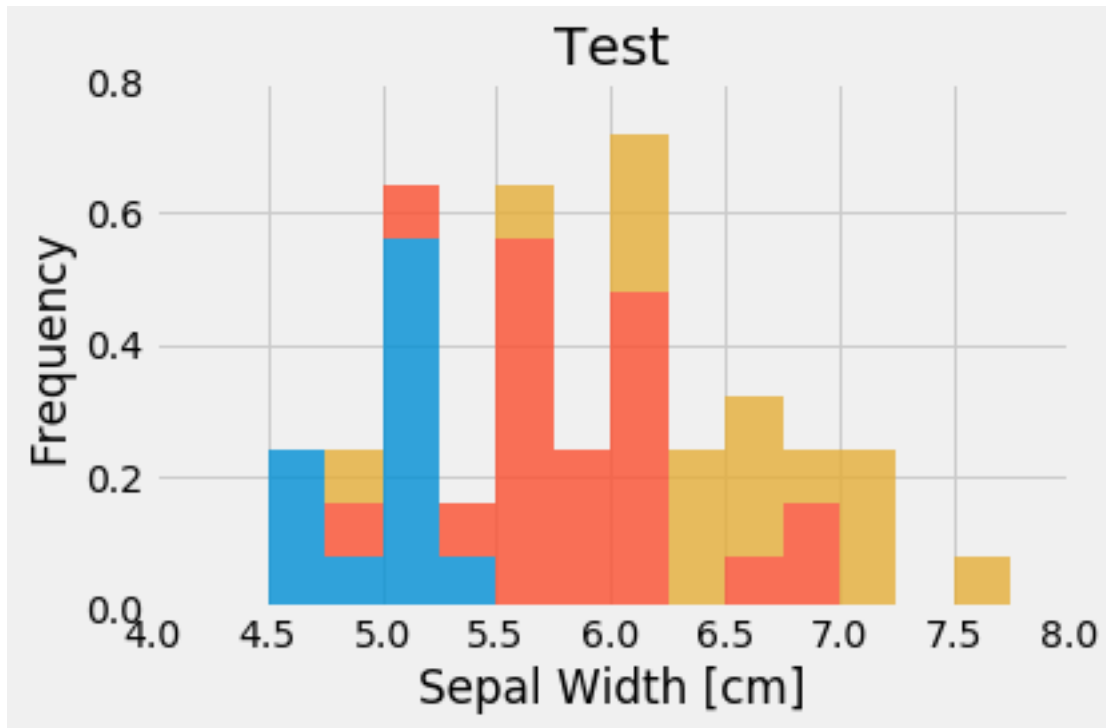
          stackhist(df[0], df['class'])
          plt.legend(['Setosa', 'Virginica', 'Versicolor'], fontsize=12)
          plt.title('All')
          plt.tight_layout()
          #plt.savefig('./all.svg')
          plt.show()

          stackhist(df_train[0], df_train['class'])
          plt.title('Train')
          plt.tight_layout()
          #plt.savefig('./train.svg')
          plt.show()

          stackhist(df_test[0], df_test['class'])
          plt.title('Test')
          plt.tight_layout()
          #plt.savefig('./test.svg')
          plt.show()

```





### 3 Holdout method and repeated sampling

```
[9]: from mlxtend.data import iris_data
      from sklearn.model_selection import train_test_split
      from sklearn.neighbors import KNeighborsClassifier
```

```
X, y = iris_data()
```

```
clf_1 = KNeighborsClassifier(n_neighbors=3,
                             weights='uniform',
                             algorithm='kd_tree',
                             leaf_size=30,
                             p=2,
                             metric='minkowski',
                             metric_params=None,
                             n_jobs=1)
```

```
[10]: rng = np.random.RandomState(seed=12345)
      seeds = np.arange(10**5)
      rng.shuffle(seeds)
      seeds = seeds[:50]
```

```
pred_2 = []
```

```

for i in seeds:
    X_train, X_test, y_train, y_test = train_test_split(X, y,
                                                         test_size=0.5,
                                                         random_state=i,
                                                         stratify=y)


    y_pred_i = clf_1.fit(X_train, y_train).predict(X_test)
    y_pred_i_acc = np.mean(y_test == y_pred_i)
    pred_2.append(y_pred_i_acc)

pred_2 = np.asarray(pred_2)
print('Average: %.2f%%' % (pred_2.mean()*100))

with plt.style.context(('fivethirtyeight')):
    plt.bar(range(0, pred_2.shape[0]), pred_2, color='gray', alpha=0.7)
    plt.axhline(pred_2.max(), color='k', linewidth=1, linestyle='--')
    plt.axhline(pred_2.min(), color='k', linewidth=1, linestyle='--')
    plt.axhspan(pred_2.min(), pred_2.max(), alpha=0.2, color='steelblue')
    plt.ylim([0, pred_2.max() + 0.1])
    plt.xlabel('Repetition')
    plt.ylabel('Accuracy')
    plt.ylim([0.5, 1.05])
    plt.tight_layout()
    #plt.savefig('figures/model-eval-iris_0.svg')
    plt.show()

```

Average: 95.41%



```
[11]: rng = np.random.RandomState(seed=12345)
seeds = np.arange(10**5)
rng.shuffle(seeds)
seeds = seeds[:50]

pred_2 = []

for i in seeds:
    X_train, X_test, y_train, y_test = train_test_split(X, y,
                                                         test_size=0.1,
                                                         random_state=i,
                                                         stratify=y)

    y_pred_i = clf_1.fit(X_train, y_train).predict(X_test)
    y_pred_i_acc = np.mean(y_test == y_pred_i)
    pred_2.append(y_pred_i_acc)

pred_2 = np.asarray(pred_2)
print('Average: %.2f%%' % (pred_2.mean()*100))

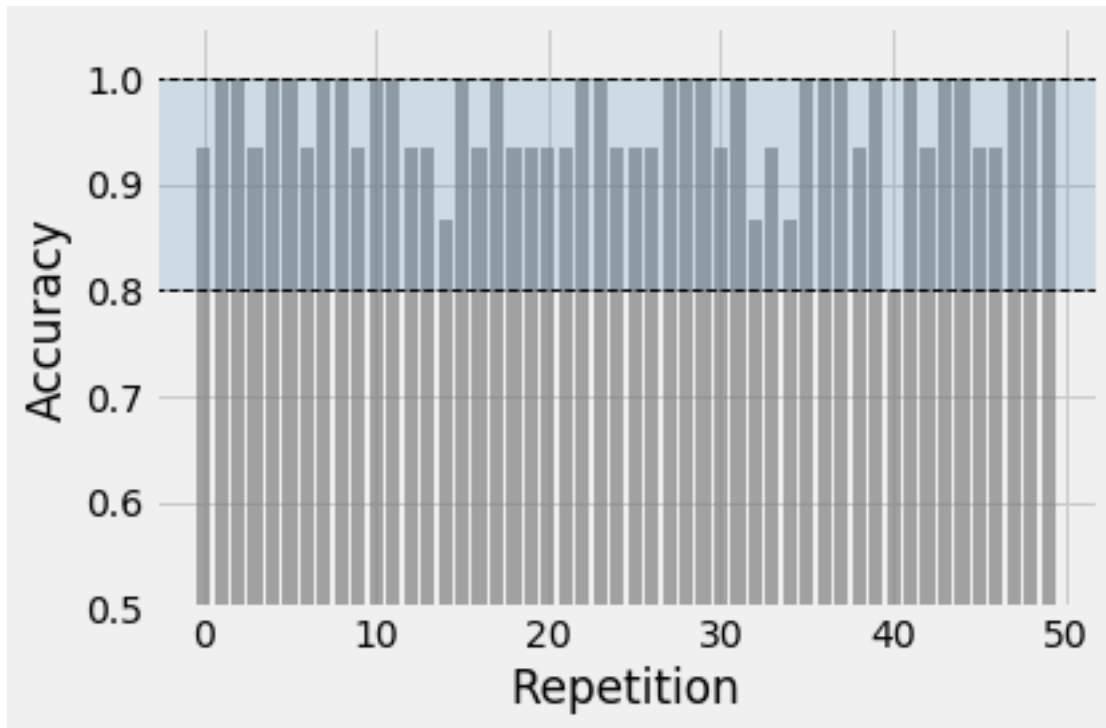
with plt.style.context(('fivethirtyeight')):
    plt.bar(range(0, pred_2.shape[0]), pred_2, color='gray', alpha=0.7)
    plt.axhline(pred_2.max(), color='k', linewidth=1, linestyle='--')
    plt.axhline(pred_2.min(), color='k', linewidth=1, linestyle='--')
    plt.axhspan(pred_2.min(), pred_2.max(), alpha=0.2, color='steelblue')
```

```

plt.ylim([0, pred_2.max() + 0.1])
plt.xlabel('Repetition')
plt.ylabel('Accuracy')
plt.ylim(0.5, 1.05)
plt.tight_layout()
#plt.savefig('figures/model-eval-iris_0_2.svg')
plt.show()

```

Average: 96.13%



### 3.1 Pessimistic Bias in Holdout

```

[12]: from mlxtend.data import mnist_data

X2, y2 = mnist_data()
X_train2, X_test2, y_train2, y_test2 = train_test_split(X2, y2,
                                                         test_size=0.3,
                                                         random_state=12,
                                                         stratify=y2)

print('Number of train examples:', y_train2.shape[0])
print('Number of test examples:', y_test2.shape[0])
print('Labels:', y_train2)

```

Number of train examples: 3500

Number of test examples: 1500  
Labels: [8 8 7 ... 6 9 8]

```
[13]: from sklearn.linear_model import LogisticRegression

clf_2 = LogisticRegression(penalty='l2',
                           dual=False,
                           tol=0.0001,
                           C=0.000001,
                           fit_intercept=True,
                           intercept_scaling=1,
                           class_weight=None,
                           random_state=12,
                           solver='lbfgs',
                           max_iter=1000,
                           multi_class='multinomial',
                           verbose=0,
                           warm_start=False,
                           n_jobs=1)


pred_train, pred_test = [], []

intervals = np.arange(500, X_train2.shape[0] + 1, 200)

for i in intervals:
    clf_2.fit(X_train2[:i], y_train2[:i])
    p_train = clf_2.score(X_train2[:i], y_train2[:i])
    p_test = clf_2.score(X_test2, y_test2)
    pred_train.append(p_train)
    pred_test.append(p_test)

with plt.style.context(('fivethirtyeight')):
    plt.plot(intervals, pred_train, marker='o', label='Train')
    plt.plot(intervals, pred_test, marker='s', label='Test')
    plt.legend(loc='best', numpoints=1)
    plt.xlim([430, X_train2.shape[0] + X_test2.shape[0]])
    plt.axvspan(X_train2.shape[0],
                X_train2.shape[0] + X_test2.shape[0],
                alpha=0.2,
                color='steelblue')
    plt.ylim([0.85, 1.0])
    plt.xlabel('Training Set Size')
    plt.ylabel('Accuracy')
    plt.tight_layout()
    #plt.savefig('figures/model-eval-mnist_0.svg')
```





```
[14]: from sklearn.ensemble import RandomForestClassifier

clf_2 = RandomForestClassifier(n_estimators=100, random_state=123)

pred_train, pred_test = [], []

intervals = np.arange(500, X_train2.shape[0] + 1, 200)


for i in intervals:
    clf_2.fit(X_train2[:i], y_train2[:i])
    p_train = clf_2.score(X_train2[:i], y_train2[:i])
    p_test = clf_2.score(X_test2, y_test2)
    pred_train.append(p_train)
    pred_test.append(p_test)

with plt.style.context(('fivethirtyeight')):
    plt.plot(intervals, pred_train, marker='o', label='Train')
    plt.plot(intervals, pred_test, marker='s', label='Test')
    plt.legend(loc='best', numpoints=1)
    plt.xlim([430, X_train2.shape[0] + X_test2.shape[0]])
    plt.axvspan(X_train2.shape[0],
                X_train2.shape[0] + X_test2.shape[0],
```

```

        alpha=0.2,
        color='steelblue')
plt.ylim([0.85, 1.0])
plt.xlabel('Training Set Size')
plt.ylabel('Accuracy')
plt.tight_layout()
#plt.savefig('figures/model-eval-mnist_0.svg')

```



```

[15]: from sklearn.neighbors import KNeighborsClassifier

clf_2 = KNeighborsClassifier(n_neighbors=3)

pred_train, pred_test = [], []

intervals = np.arange(500, X_train2.shape[0] + 1, 200)


for i in intervals:
    clf_2.fit(X_train2[:i], y_train2[:i])
    p_train = clf_2.score(X_train2[:i], y_train2[:i])
    p_test = clf_2.score(X_test2, y_test2)
    pred_train.append(p_train)
    pred_test.append(p_test)

```

```

with plt.style.context(('fivethirtyeight')):
    plt.plot(intervals, pred_train, marker='o', label='Train')
    plt.plot(intervals, pred_test, marker='s', label='Test')
    plt.legend(loc='best', numpoints=1)
    plt.xlim([430, X_train2.shape[0] + X_test2.shape[0]])
    plt.axvspan(X_train2.shape[0],
                X_train2.shape[0] + X_test2.shape[0],
                alpha=0.2,
                color='steelblue')
    plt.ylim([0.85, 1.0])
    plt.xlabel('Training Set Size')
    plt.ylabel('Accuracy')
    plt.tight_layout()
    #plt.savefig('figures/model-eval-mnist_0.svg')

```



## 3.2 Bootstrapping

### 3.2.1 Iris

```

[16]: from sklearn.neighbors import KNeighborsClassifier

rng = np.random.RandomState(seed=12345)

```

```

X, y = iris_data()

clf = KNeighborsClassifier(n_neighbors=3,
                           weights='uniform',
                           algorithm='kd_tree',
                           leaf_size=30,
                           p=2,
                           metric='minkowski',
                           metric_params=None,
                           n_jobs=1)

idx = np.arange(y.shape[0])

accuracies = []

for i in range(200):

    train_idx = rng.choice(idx, size=idx.shape[0], replace=True)
    test_idx = np.setdiff1d(idx, train_idx, assume_unique=False)

    boot_train_X, boot_train_y = X[train_idx], y[train_idx]
    boot_test_X, boot_test_y = X[test_idx], y[test_idx]

    clf.fit(boot_train_X, boot_train_y)
    acc = clf.score(boot_test_X, boot_test_y)
    accuracies.append(acc)

```

```

[17]: mean = np.mean(accuracies)

#se = np.sqrt( (1. / (100-1)) * np.sum([(acc - mean)**2 for acc in
    ↪ accuracies]))
#ci = 1.984 * se

se = np.sqrt( (1. / (200-1)) * np.sum([(acc - mean)**2 for acc in accuracies]))
ci = 1.97 * se

lower = np.percentile(accuracies, 2.5)
upper = np.percentile(accuracies, 97.5)

fig, ax = plt.subplots(figsize=(8, 4))
ax.vlines(mean, [0], 80, lw=2.5, linestyle='-', label='mean')
#ax.vlines(med, [0], 60, lw=2.5, linestyle='--', label='median')
ax.vlines(lower, [0], 15, lw=2.5, linestyle='-.', label='CI95 percentile')
ax.vlines(upper, [0], 15, lw=2.5, linestyle='-.')

ax.vlines(mean + ci, [0], 15, lw=2.5, linestyle=':', label='CI95 standard')

```

```


ax.vlines(mean - ci, [0], 15, lw=2.5, linestyle=':')

ax.hist(accuracies, bins=7,
        color='#0080ff', edgecolor="none",
        alpha=0.3)
plt.legend(loc='upper left')

plt.xlim([0.86, 1.01])
plt.tight_layout()
#plt.savefig('figures/bootstrap-histo-1.svg')

plt.show()

```



### 3.2.2 5k MNIST Subset

```

[18]: from mlxtend.data import mnist_data
      from sklearn.linear_model import LogisticRegression

rng = np.random.RandomState(seed=12345)

X, y = mnist_data()

clf = LogisticRegression(penalty='l2',
                        dual=False,
                        tol=0.0001,
                        C=0.000001,

```

```

        fit_intercept=True,
        intercept_scaling=1,
        class_weight=None,
        random_state=12,
        solver='lbfgs',
        max_iter=1000,
        multi_class='multinomial',
        verbose=0,
        warm_start=False,
        n_jobs=1)

idx = np.arange(y.shape[0])

accuracies = []

for i in range(200):

    train_idx = rng.choice(idx, size=idx.shape[0], replace=True)
    test_idx = np.setdiff1d(idx, train_idx, assume_unique=False)

    boot_train_X, boot_train_y = X[train_idx], y[train_idx]
    boot_test_X, boot_test_y = X[test_idx], y[test_idx]

    clf.fit(boot_train_X, boot_train_y)
    acc = clf.score(boot_test_X, boot_test_y)
    accuracies.append(acc)

```

```

[19]: mean = np.mean(accuracies)

#se = np.sqrt( (1. / (100-1)) * np.sum([(acc - mean)**2 for acc in
    ↪ accuracies]))
#ci = 1.984 * se

se = np.sqrt( (1. / (200-1)) * np.sum([(acc - mean)**2 for acc in accuracies]))
ci = 1.97 * se

lower = np.percentile(accuracies, 2.5)
upper = np.percentile(accuracies, 97.5)

fig, ax = plt.subplots(figsize=(8, 4))
ax.vlines(mean, [0], 40, lw=2.5, linestyle='-', label='mean')
#ax.vlines(med, [0], 60, lw=2.5, linestyle='--', label='median')
ax.vlines(lower, [0], 15, lw=2.5, linestyle='-.', label='CI95 percentile')
ax.vlines(upper, [0], 15, lw=2.5, linestyle='-.')

ax.vlines(mean + ci, [0], 15, lw=2.5, linestyle=':', label='CI95 standard')

```

```

ax.vlines(mean - ci, [0], 15, lw=2.5, linestyle=':')


ax.hist(accuracies, bins=11,
        color='#0080ff', edgecolor="none",
        alpha=0.3)
plt.legend(loc='upper left')

plt.xlim([0.885, 0.922])
plt.tight_layout()

# plt.savefig('figures/bootstrap-histo-2.svg')

plt.show()

```



### 3.2.3 Out of Bag Bootstrap

```

[20]: from mlxtend.evaluate import BootstrapOutOfBag
import numpy as np

oob = BootstrapOutOfBag(n_splits=3)
for train, test in oob.split(np.array([1, 2, 3, 4, 5])):
    print(train, test)

```

```

[1 1 4 1 1] [0 2 3]
[0 1 1 3 0] [2 4]
[3 3 3 3 0] [1 2 4]

```

```
[21]: from mlxtend.data import iris_data
      from sklearn.tree import DecisionTreeClassifier
      from sklearn.model_selection import cross_val_score

      X, y = iris_data()
      model = DecisionTreeClassifier()
```

```
[22]: cv_scores = cross_val_score(model, X, y, cv=5)
      print('CV scores', cv_scores)
      print('Mean CV score', np.mean(cv_scores))
      print('CV score Std', np.std(cv_scores))
```

```
CV scores [0.96666667 0.96666667 0.9          1.          1.          ]
Mean CV score 0.9666666666666668
CV score Std 0.036514837167011066
```

```
[23]: bootstrap_scores = \
      cross_val_score(model, X, y,
                      cv=BootstrapOutOfBag(n_splits=5, random_seed=456))

      print('Bootstrap scores', bootstrap_scores)
      print('Mean Bootstrap score', np.mean(bootstrap_scores))
      print('Score Std', np.std(bootstrap_scores))
```

```
Bootstrap scores [0.92727273 0.98113208 0.92592593 0.98305085 0.90909091]
Mean Bootstrap score 0.9452944970437775
Score Std 0.030726208804992074
```

```
[24]: bootstrap_scores = \
      cross_val_score(model, X, y,
                      cv=BootstrapOutOfBag(n_splits=200, random_seed=456))

      print('Mean Bootstrap score', np.mean(bootstrap_scores))
      print('Score Std', np.std(bootstrap_scores))
```

```
Mean Bootstrap score 0.9448327329412445
Score Std 0.030836791007487842
```

```
[25]: lower = np.percentile(bootstrap_scores, 2.5)
      upper = np.percentile(bootstrap_scores, 97.5)
      print('95%% Confidence interval: [%.2f, %.2f]' % (100*lower, 100*upper))
```

```
95% Confidence interval: [87.27, 100.00]
```



### 3.2.4 .632 Bootstrap

```
[26]: from mlxtend.evaluate import bootstrap_point632_score
```

```
bootstrap_scores = bootstrap_point632_score(model, X, y, n_splits=200)
print('Mean Bootstrap score', np.mean(bootstrap_scores))
print('Score Std', np.std(bootstrap_scores))
```

Mean Bootstrap score 0.9627083103060582  
Score Std 0.016892377745312225

```
[27]: lower = np.percentile(bootstrap_scores, 2.5)
upper = np.percentile(bootstrap_scores, 97.5)
print('95%% Confidence interval: [%.2f, %.2f]' % (100*lower, 100*upper))
```

95% Confidence interval: [92.81, 98.89]